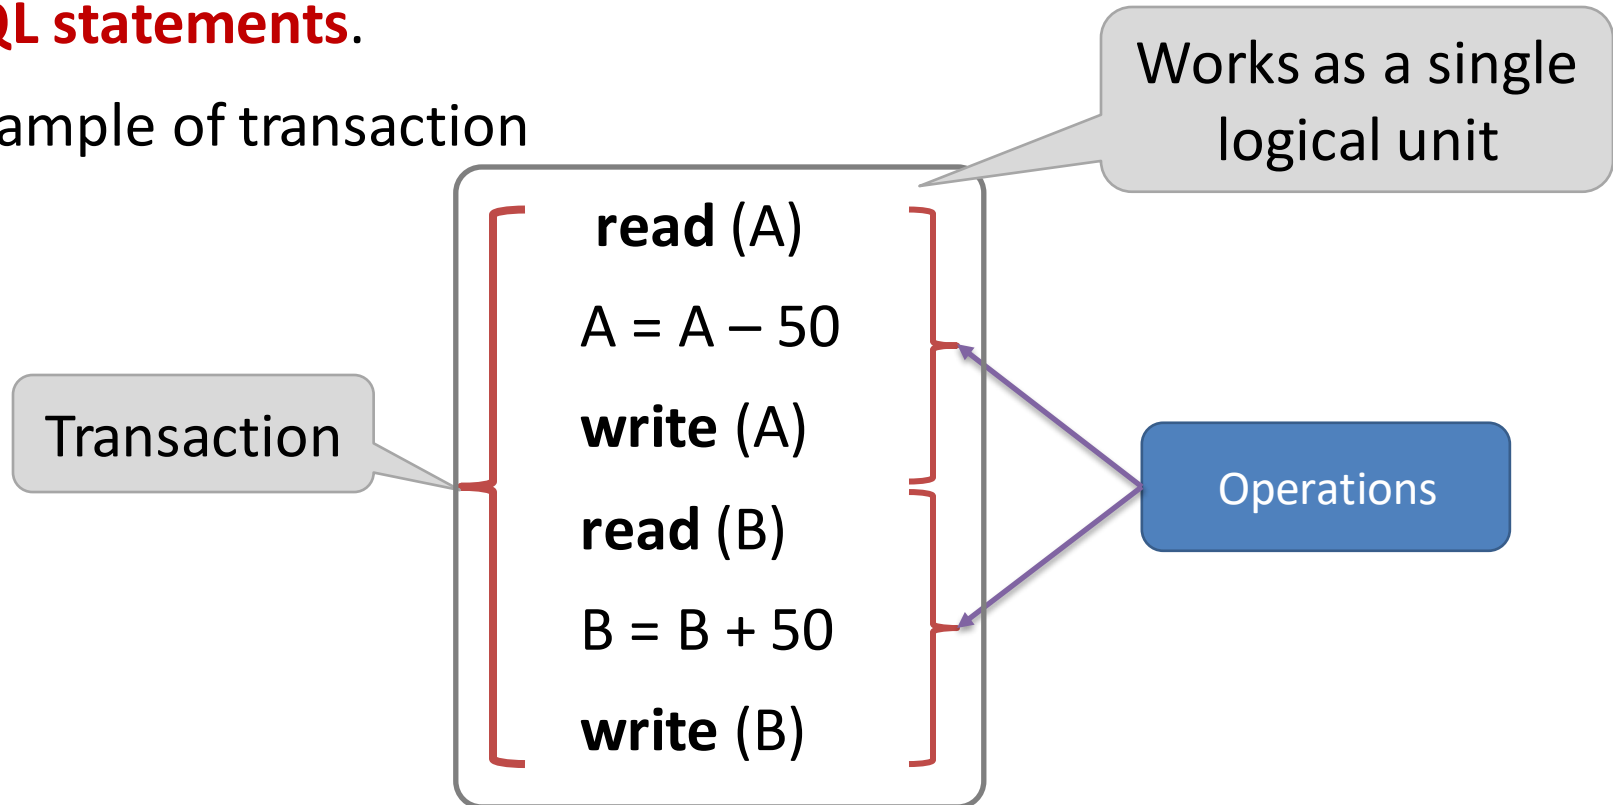


Transaction Management

BY: PROF. PRACHI SHAH

What is transaction?

- A transaction is a **sequence of operations performed as a single logical unit of work.**
- A transaction is a **logical unit of work that contains one or more SQL statements.**
- Example of transaction



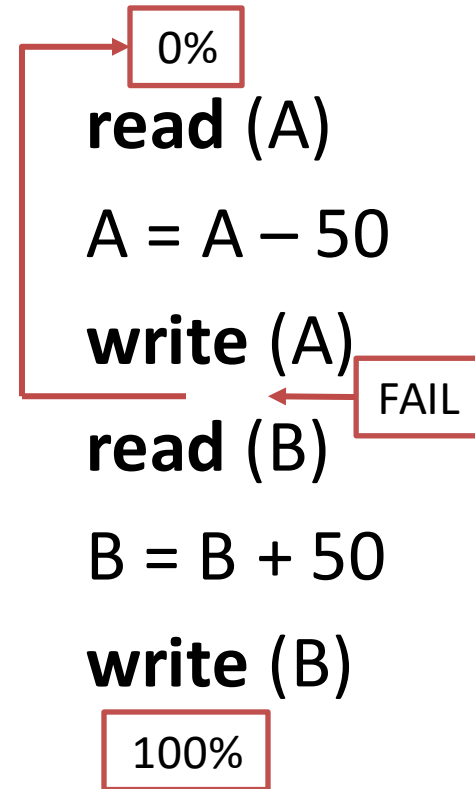
ACID properties of transaction

- **Atomicity** (Either transaction execute 0% or 100%)
- **Consistency** (database must remain in a consistent state after any transaction)
- **Isolation** (Multiple transaction may execute concurrently, but each transaction should feel like its executing in isolation)
- **Durability** (Once a transaction completed successfully, the changes it has made into the database should be permanent)

ACID properties of transaction

- **Atomicity**

- This property states that a **transaction must be treated as an atomic unit**, that is, **either all of its operations are executed or none.**
- **Either transaction execute 0% or 100%.**
- For example, consider a transaction to transfer Rs. 50 from account A to account B.
- In this transaction, if Rs. 50 is deducted from account A then it must be added to account B.



ACID properties of transaction

- **Consistency**

- The **database must remain in a consistent state after any transaction.**
- If the database was in a consistent state before the execution of a transaction, it must remain consistent after the execution of the transaction as well.
- In our example, total of A and B must remain same before and after the execution of transaction.

A=500, B=500

A+B=1000

read (A)

A = A - 50

write (A)

read (B)

B = B + 50

write (B)

A=450, B=550

A+B=1000

ACID properties of transaction

- Isolation

- Changes occurring in a particular transaction will not be visible to any other transaction until it has been committed.
- Intermediate transaction results must be hidden from other concurrently executed transactions.
- In our example once our transaction starts from first step (step 1) its result should not be access by any other transaction until last step (step 6) is completed.

read (A)

$A = A - 50$

write (A)

read (B)

$B = B + 50$

write (B)

ACID properties of transaction

- **Durability**

- After a transaction completes successfully, **the changes it has made to the database persist (permanent)**, even if there are system failures.
- Once our transaction completed up to last step (step 6) its result must be stored permanently. It should not be removed if system fails.

A=500, B=500

read (A)

$A = A - 50$

write (A)

read (B)


$B = B + 50$

write (B)

A=450, B=550

Transaction State Diagram \ State Transition Diagram

- When a transaction executes its final operation, it is said to be in a partially committed state.




- This is the initial state.
- The transaction stays in this state while it is executing.

read (A)
 $A = A - 50$
write (A)
read (B)
 $B = B + 50$
write (B)
Commit

Transaction State Diagram \ State Transition Diagram

- The transaction enters in this state after successful completion of the transaction.
- We cannot abort or rollback a committed transaction.



- Discover that normal execution can no longer proceed.
- Once a transaction cannot be completed, any changes that it made must be undone rolling it back.

read (A)
 $A = A - 50$
write (A)
read (B)
 $B = B + 50$
write (B)
Commit

Transaction State Diagram \ State Transition Diagram

- Active
 - This is the **initial state**.
 - The transaction **stays in this state while it is executing**.
- Partial Committed
 - When a transaction **executes its final operation/ instruction**, it is said to be in a partially committed state.
- Failed
 - Discover that **normal execution can no longer proceed**.
 - Once a **transaction cannot be completed**, any **changes that it made must be undone rolling it back**.

Transaction State Diagram \ State Transition Diagram

- Committed

- The transaction enters in this state after **successful completion of the transaction (after committing transaction)**.
- We **cannot abort or rollback a committed transaction**.

- Aborted

- The **state after the transaction has been rolled back** and the **database has been restored to its state prior to the start of the transaction**.

What is schedule?

- A schedule is a **process of grouping the transactions** into one and **executing them in a predefined order**.
- A schedule is the **chronological (sequential) order** in which **instructions are executed** in a system.
- A schedule is required in a database because **when some transactions execute in parallel, they may affect the result of the transaction**.
- Means if one transaction is updating the values which the other transaction is accessing, then the order of these two transactions will change the result of another transaction.
- Hence a schedule is created to execute the transactions.

Example of schedule

T1	T2	A=B=1000
Read (A) A = A - 50 Write (A) Read (B) B = B + 50 Write (B) Commit	Read (A) temp = A * 0.1 A = A - temp Write (A) Read (B) B = B + temp Write (B) Commit	Read (1000) A = 1000 - 50 Write (950) Read (1000) B = 1000 + 50 Write (1050) Commit Read (950) temp = 950 * 0.1 A = 950 - 95 Write (855) Read (1050) B = 1050 + 95 Write (1145) Commit

Example of schedule

T1	T2	A=B=1000
<p>Read (A) Temp = A * 0.1 A = A - temp Write (A) Read (B) B = B + temp Write (B) Commit</p>	<p>Read (A) A = A - 50 Write (A) Read (B) B = B + 50 Write (B) Commit</p>	<p>Read (1000) Temp = 1000 * 0.1 A = 1000 - 100 Write (900) Read (1000) B = 1000 + 100 Write (1100) Commit Read (900) A = 900 - 50 Write (850) Read (1100) B = 1100 + 50 Write (1150) Commit</p>

Serial schedule

- A serial schedule is one in which **no transaction starts until a running transaction has ended.**
- Transactions are executed one after the other.
- This type of schedule is called a serial schedule, as transactions are executed in a serial manner.

Example of serial schedule

T1	T2	A=B=1000
Read (A) Temp = A * 0.1 A = A - temp Write (A) Read (B) B = B + temp Write (B) Commit	Read (A) A = A - 50 Write (A) Read (B) B = B + 50 Write (B) Commit	Read (1000) Temp = 1000 * 0.1 A = 1000 - 100 Write (900) Read (1000) B = 1000 + 100 Write (1100) Commit Read (900) A = 900 - 50 Write (850) Read (1100) B = 1100 + 50 Write (1150) Commit

Example of serial schedule

T1	T2	A=B=1000
Read (A) Temp = A * 0.1 A = A - temp Write (A) Read (B) B = B + temp Write (B) Commit	Read (A) A = A - 50 Write (A) Read (B) B = B + 50 Write (B) Commit	Read (1000) A = 1000 - 50 Write (950) Read (1000) B = 1000 + 50 Write (1050) Commit Read (950) Temp = 950 * 0.1 A = 950 - temp (95) Write (855) Read (1050) B = 1050 + temp (95) Write (1145) Commit

Concurrent / Interleaved schedule

- Schedule that **interleave the execution of different transactions.**
- Means **second transaction is started before the first one could end** and **execution can switch between the transactions back and forth.**

Example of interleaved schedule

T1	T2
Read (A) A = A - 50 Write (A)	Read (A) Temp = A * 0.1 A = A - temp Write (A)
Read (B) B = B + 50 Write (B) Commit	Read (B) B = B + temp Write (B) Commit

Example of interleaved schedule

T1	T2	A=B=1000
Read (A) A = A - 50		Read (1000) A = 1000 - 50
	Read (A) Temp = A * 0.1 A = A - temp Write (A)	Read (1000) Temp = 1000 * 0.1 A = 1000 - temp (100)
Write (A) Read (B) B = B + 50 Write (B) Commit		Write (900) Write (950) Read (1000) B = 1000 + 50
	Read (B) B = B + temp Write (B) Commit	Write (1050) Commit Read (1050) B = 1050 + temp (100)
		Write (1150) Commit

Equivalent schedule

- If two schedules **produce the same result after execution**, they are said to be equivalent schedule.
- They may yield the same result for some value and different results for another set of values.
- That's why this equivalence is not generally considered significant.

Equivalent schedule

T1	T2
Read (A) $A = A - 50$ Write (A) Commit	Read (A) $\text{Temp} = A * 0.1$ $A = A - \text{temp}$ Write (A) Commit
Read (B) $B = B + 50$ Write (B) Commit	Read (B) $B = B + \text{temp}$ Write (B) Commit

Both schedules are equivalent

In both schedules the sum " $A + B$ " is preserved.

T1	T2
Read (A) $A = A - 50$ Write (A) Read (B) $B = B + 50$ Write (B) Commit	Read (A) $\text{Temp} = A * 0.1$ $A = A - \text{temp}$ Write (A) Read (B) $B = B + \text{temp}$ Write (B) Commit

Serializability

- Serializability is a concept which will help us to identify valid concurrent schedules.
- A schedule is serializable **if it is equivalent to a serial schedule**.
- In serial schedules, **only one transaction is allowed to execute at a time** i.e. no concurrency is allowed.
- Whereas in serializable schedules, **multiple transactions can execute simultaneously** i.e. concurrency is allowed.
- Types (forms) of serializability
 1. Conflict serializability
 2. View serializability

Conflicting instructions

- Let I_i and I_j be two instructions of transactions T_i and T_j respectively.

1. $I_i = \text{read}(Q)$, $I_j = \text{read}(Q)$

I_i and I_j **don't conflict**

T1	T2
read (Q)	
	read (Q)

T1	T2
	read (Q)
read (Q)	

2. $I_i = \text{read}(Q)$, $I_j = \text{write}(Q)$

I_i and I_j **conflict**

T1	T2
read (Q)	
	write(Q)

T1	T2
	write(Q)
read (Q)	

3. $I_i = \text{write}(Q)$, $I_j = \text{read}(Q)$

I_i and I_j **conflict**

T1	T2
write(Q)	
	read (Q)

T1	T2
	read (Q)
write(Q)	

4. $I_i = \text{write}(Q)$, $I_j = \text{write}(Q)$

I_i and I_j **conflict**

T1	T2
write(Q)	
	write(Q)

T1	T2
	write(Q)
write(Q)	

Conflicting instructions

- Instructions I_i and I_j **conflict** if and only if there exists some item Q **accessed by both I_i and I_j , and at least one of these instructions write Q .**
- If **both the transactions access different data item** then they are **not conflict**.

Conflict serializability

- If a given **schedule can be converted into a serial schedule by swapping its non-conflicting operations**, then it is called as a conflict serializable schedule.

Conflict serializability (example)

T1	T2
Read (A) A = A - 50 Write (A)	<div>Read (A) Temp = A * 0.1 A = A - temp Write (A)</div> <div>Read (B) B = B + 50 Write (B) Commit</div>
	Read (B) B = B + temp Write (B) Commit

T1	T2
Read (A) A = A - 50 Write (A) Read (B) B = B + 50 Write (B) Commit	<div>Read (A) Temp = A * 0.1 A = A - temp Write (A) Read (B) B = B + temp Write (B) Commit</div>

Conflict serializability

- Example of a schedule that is not conflict serializable:

T1	T2
Read (A)	Write (A)
Write (A)	

- We are **unable to swap instructions** in the above schedule to obtain either the serial schedule $\langle T1, T2 \rangle$, or the serial schedule $\langle T2, T1 \rangle$.

Key points Serializability


- Every conflict serializable is also a view serializable schedule (NOT vice versa)

- **Testing of Conflict serializability : Using precedence graph.**
 1. Create number of nodes = number of transactions
 2. Check for conflict pairs (R-W, W-R, W-W on same data item) and draw the edges between them.
 3. Check for LOOP/CYCLE in graph. If loop is there, NOT CONFLICT SERIALIZABLE.

Example 1: Conflict Serializability

Check whether the given schedule S is conflict serializable or not-

$S : R_1(A) , R_2(A) , R_1(B) , R_2(B) , R_3(B) , W_1(A) , W_2(B)$



T1	T2	T3
$R_1(A)$		
	$R_2(A)$	
$R_1(B)$		
	$R_2(B)$	
		$R_3(B)$
$W_1(A)$		
	$W_2(B)$	

S is not conflict serializable.

Example 2: Conflict Serializability

- Consider a schedule S with two transactions T_1 and T_2 as follows;
S: $R_1(x); W_2(x); R_1(x); W_1(y); \text{commit}_1; \text{commit}_2$;
- Is the schedule S conflict serializable?

Instruction	T_1	T_2
1	$R(x)$	$W(x)$
2		
3	$R(x)$	
4	$W(y)$	$Commit$
5	$Commit$	
6		

Example 2: Conflict Serializability

- Consider a schedule S with two transactions T_1 and T_2 as follows;
 $S: R_1(x); W_2(x); R_1(x); W_1(y); \text{commit}_1; \text{commit}_2;$
- Is the schedule S conflict serializable?



Schedule S is not conflict serializable schedule.


Example 3: Conflict Serializability

Check whether the given schedule S is conflict serializable or not-

T1	T2	T3	T4
	R(X)		
		W(X) Commit	
W(X) Commit			
	W(Y) R(Z) Commit		
			R(X) R(Y) Commit

Example 3: Conflict Serializability

Check whether the given schedule S is conflict serializable or not-



**S is conflict serializable.
Serial Order:**

T2-T3-T1-T4

View serializability

- Let S_1 and S_2 be two schedules with the same set of transactions. S_1 and S_2 are **view equivalent if the following three conditions** are satisfied, for each data item Q
 1. Initial Read
 2. Updated Read
 3. Final Write

Initial Read

- If in **schedule S1**, transaction T_i reads the initial value of Q , then in **schedule S2** also transaction T_i must read the initial value of Q .

S1		S3		S2	
T1	T2	T1	T2	T1	T2
Read (A)	Write (A)	Write (A)	Read (A)	Read (A)	Write (A)

- Above two schedules S1 and S3 are **not view equivalent** because initial read operation in S1 is done by T1 and in S3 it is done by T2.
- Above two schedules S1 and S2 **are view equivalent** because initial read operation in S1 is done by T1 and in S2 it is also done by T1.

Updated Read

- If in **schedule S1** transaction T_i executes **read(Q)**, and that value was produced by transaction T_j (if any), then in **schedule S2** also transaction T_i must read the value of Q that was produced by transaction T_j .

S1		
T1	T2	T3
Write (A)	Write (A)	Read (A)

S3		
T1	T2	T3
Write (A)	Write (A)	Read (A)

- Above two schedules are **not view serializable** because,
- in S1, T3 is reading A that is updated by T2 and
- in S3, T3 is reading A which is updated by T1.

Updated Read

- If in **schedule S1** transaction T_i executes **read(Q)**, and that value was produced by transaction T_j (if any), then in **schedule S2** also transaction T_i must read the value of Q that was produced by the same **write(Q)** operation of transaction T_j .

S1		
T1	T2	T3
Write (A)	Read (A) Write (A)	Read (A)

S2		
T1	T2	T3
Write (A)	Read (A) Write (A)	Read (A)

- Above two schedules are view equal because,
- in S1, T3 is reading A that is updated by T2 and
- in S3 also, T3 is reading A which is updated by T2.

Final Write

- If T_i performs the final write on the data value in $S1$, then T_i also performs the final write on the data value in $S2$.

S1		
T1	T2	T3
Write (A)	Read (A)	Write (A)

S2		
T1	T2	T3
Write (A)	Read (A)	Write (A)

- Above two schedules is view equal because final write operation
- in $S1$ is done by $T3$ and
- in $S2$ also the final write operation is also done by $T3$.

View Serializability

T1	T2	T3
$R_1(A)$		
	$W_2(A)$	
$W_1(A)$		
		$W_3(A)$

Possible serial schedules?

T1,T2,T3

T1,T3,T2

T2,T1,T3

T2,T3,T1

T3,T1,T2

T3,T2,T1

You can compare it with any serial schedule, **and if all three conditions are true, It is view serializable.**

	T1	T2	T3
1. Initial read?	$R_1(A)$		
2. Updated read?	$W_1(A)$		
3. Final write?		$W_2(A)$	$W_3(A)$

Key points Serializability

- *Every conflict serializable is also a view serializable schedule (NOT vice versa)*

Example 4: View Serializability


Check whether the given schedule S is view serializable or not-

T1	T2	T3
R (A)		
	R (A)	
		W (A)
W (A)		

**Is it conflict
serializable?**

Example 4: View Serializability

Check whether the given schedule S is view serializable or not-



Is it conflict
serializable?
NO

Is it View
serializable?

Example 4: View Serializability

Check whether the given schedule S is view serializable or not-

Is it View
serializable?
NO

Example 5: View Serializability

Check whether the given schedule S is view serializable or not. If yes, then give the serial schedule.

S : $R_1(A)$, $W_2(A)$, $R_3(A)$, $W_1(A)$, $W_3(A)$

Recoverable Schedule

- **Recoverable schedule:**
- A schedule S is recoverable if
 1. Transaction T_i have read(x), that has been previously written by some other transaction T_j .
 2. Transaction T_j must commit before T_i commits.
- Suppose transaction T_j fails/aborts, then we have to rollback all operations done by T_j , *and it is also necessary to abort T_i (which is dependent on T_j).*
- Order of Commit/Abort Matters.

if T_9 commits after
read(A)

T_8	T_9
read(A) write(A) read(B)	 read(A)

Recoverable Schedule

- S_a is not recoverable:
 - $S_a: r_1(X); w_1(X); r_2(X); r_1(Y); w_2(X); c_2; w_1(Y); c_1;$
- Consider the following schedules:
 - $S_c: r_1(X); w_1(X); r_2(X); r_1(Y); w_2(X); c_2; a_1;$
 - $S_d: r_1(X); w_1(X); r_2(X); r_1(Y); w_2(X); w_1(Y); c_1; c_2;$
 - $S_e: r_1(X); w_1(X); r_2(X); r_1(Y); w_2(X); w_1(Y); a_1; a_2;$
- S_c is **not recoverable** because:
- S_d is **recoverable** because:
- S_e is **recoverable** because:

Cascade less Schedule

- Cascading Rollbacks / aborts

T_{10}	T_{11}	T_{12}
read(A) read(B) write(A)	read(A) write(A)	read(A)

- Cascade less schedule : **Eliminates Cascading rollbacks:**
 1. T_i read(x), Previously written by T_j
 2. T_j must commit before T_i reads.
- Every cascade less schedule is also recoverable.

Cascade less Schedule

T1	T2
R(X)	
W(X)	
	R(Y)
COMMIT;	
	R(X)


CASCADE LESS?

RECOVERABLE?

Strict Schedule

- **Strict schedule:**
- A schedule in which a transaction can **neither read or write** an item X until the last transaction that **wrote** X has committed.
- S_a is not strict:
 - $S_a: w1(X, 5); w2(X, 8); a1$
- Suppose the value of X was originally 9.
- If T1 aborts, as in S_a , the recovery system will restore the value of X to 9, even though it has already been changed to 8 by T2, thus leading to incorrect results.
- Although S_a is cascade-less, it is not strict
 - It permits T2 to write X even though T1 that last wrote X had not yet committed (or aborted).

Schedules



GATE Example Schedule

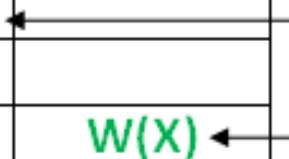
- Consider the following database schedule with two transactions, T1 and T2.

$S = r_2(X); r_1(X); r_2(Y); w_1(X); r_1(Y); w_2(X); a_1; a_2;$

- S is recoverable?
- S is cascade less?
- S is strict?

GATE Example Schedule

T1	T2
	R(X)
R(X)	
	R(Y)
W(X)	
R(Y)	
	W(X)
a1	
	a2



- Given schedule is recoverable.
- Given schedule is cascade less.
- Given schedule is NOT strict.

The End..!!
Thank You....